

Data Set Information:

This file is part of APS Failure and Operational Data for Scania Trucks.

Copyright (c) <2016>

This program (APS Failure and Operational Data for Scania Trucks) is free software: you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation, either version 3 of the License, or (at your option) any later version.

This program is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

You should have received a copy of the GNU General Public License along with this program. If not, see <[\[Web Link\]](#)>.

-
1. Title: APS Failure at Scania Trucks
 2. Source Information

-- Creator: Scania CV AB Vagnmakarvägen 1 151 32 Södertälje Stockholm Sweden -- Donor: Tony Lindgren (tony '@' dsv.su.se) and Jonas Biteus (jonas.biteus '@' scania.com) -- Date: September, 2016

3. Past Usage: Industrial Challenge 2016 at The 15th International Symposium on Intelligent Data Analysis (IDA)

-- Results: The top three contestants | Score | Number of Type 1 faults | Number of Type 2 faults

Camila F. Costa and Mario A. Nascimento | 9920 | 542 | 9 Christopher Gondek, Daniel Hafner and Oliver R. Sampson | 10900 | 490 | 12 Sumeet Garnaik, Sushovan Das, Rama Syamala Sreepada and Bidyut Kr. Patra | 11480 | 398 | 15

4. Relevant Information:

-- Introduction The dataset consists of data collected from heavy Scania trucks in everyday usage. The system in focus is the Air Pressure system (APS) which generates pressurised air that are utilized in various functions in a truck, such as braking and gear changes. The datasets' positive class consists of component failures for a specific component of the APS system. The negative class consists of trucks with failures for components not related to the APS. The data consists of a subset of all available data, selected by experts.

-- Challenge metric

Cost-metric of miss-classification:

Predicted class | True class || pos | neg |

pos | - | Cost_1 |

neg | Cost_2 | - |

Cost_1 = 10 and cost_2 = 500

The total cost of a prediction model the sum of 'Cost_1' multiplied by the number of Instances with type 1 failure and 'Cost_2' with the number of instances with type 2 failure, resulting in a 'Total_cost'.

In this case Cost_1 refers to the cost that an unnessecary check needs to be done by an mechanic at an workshop, while Cost_2 refer to the cost of missing a faulty truck, which may cause a breakdown.

Total_cost = Cost_1No_Instances + Cost_2No_Instances.

5. Number of Instances: The training set contains 60000 examples in total in which 59000 belong to the negative class and 1000 positive class. The test set contains 16000 examples.
6. Number of Attributes: 171
7. Attribute Information: The attribute names of the data have been anonymized for proprietary reasons. It consists of both single numerical counters and histograms consisting of bins with different conditions. Typically the histograms have open-ended conditions at each end. For example if we measuring the ambient temperature "T" then the histogram could be defined with 4 bins where:

bin 1 collect values for temperature T < -20 bin 2 collect values for temperature T >= -20 and T < 0 bin 3 collect values for temperature T >= 0 and T < 20 bin 4 collect values for temperature T > 20

| b1 | b2 | b3 | b4 |

-20 0 20

The attributes are as follows: class, then anonymized operational data. The operational data have an identifier and a bin id, like 'Identifier_Bin'. In total there are 171 attributes, of which 7 are histogram variabels. Missing values are denoted by 'na'.

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import warnings
import seaborn as sns
warnings.filterwarnings('ignore')
%matplotlib inline
```

```
In [2]: df1=pd.read_csv(r"C:\Users\Mukul\Downloads\aps_failure_training_set.csv")
```

```
In [3]: df2=pd.read_csv(r"C:\Users\Mukul\Downloads\aps_failure_test_set.csv")
```

```
In [4]: df=pd.concat([df1,df2])
```

```
In [8]: ### joining both dataset and resetting index
df=pd.concat([df1, df2])

df.reset_index(inplace=True)
## dropping index feature as it is not required
df.drop('index',axis=1,inplace=True)

# replacing na values with NaN Values
df=df.replace('na' , np.NaN)

df.head()
```

Out[8]:

```
      class aa_000 ab_000      ac_000 ad_000 ae_000 af_000 ag_000 ag_001 ag_002 ... ee_002 ee_003 ee_004 ee_005 ee_006 ee_007 ee_008 ee_00
0   neg  76698    NaN  2130706438     280     0     0     0     0     0 ... 1240520  493384  721044  469792  339156  157956  73224 ...
1   neg  33058    NaN        0    NaN     0     0     0     0     0 ... 421400  178064  293306  245416  133654  81140   97576  150 ...
2   neg  41040    NaN        228    100     0     0     0     0     0 ... 277378  159812  423992  409564  320746  158022  95128  51 ...
3   neg     12     0       70     66     0    10     0     0     0 ... 240      46      58      44      10      0      0      0 ...
4   neg  60874    NaN     1368    458     0     0     0     0     0 ... 622012  229790  405298  347188  286954  311560  433954  121 ...
```

5 rows × 171 columns

```
In [11]: df.tail()
```

Out[11]:

```
      class aa_000 ab_000      ac_000 ad_000 ae_000 af_000 ag_000 ag_001 ag_002 ... ee_002 ee_003 ee_004 ee_005 ee_006 ee_007 ee_008 ee_00
75995   neg  81852    NaN  2130706432     892     0     0     0     0 ... 632658  273242  510354  373918  349840  317840  960024  2 ...
75996   neg     18     0      52     46     8    26     0     0     0 ... 266      44      46      14      2      0      0      0 ...
75997   neg  79636    NaN     1670    1518     0     0     0     0     0 ... 806832  449962  778826  581558  375498  222866  358934  1 ...
75998   neg    110    NaN     36     32     0     0     0     0     0 ... 588      210     180     544    1004    1338     74 ...
75999   neg      8     0       6     4     2     2     0     0     0 ... 46      10      48      14      42      46      0      0 ...
```

5 rows × 171 columns

```
In [12]: df.shape
```

Out[12]: (76000, 171)

```
In [13]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 76000 entries, 0 to 75999
Columns: 171 entries, class to eg_000
dtypes: int64(1), object(170)
memory usage: 99.2+ MB
```

```
In [14]: # check the null values  
df.isnull().sum()
```

```
Out[14]: class      0  
aa_000      0  
ab_000    58692  
ac_000     4261  
ad_000    18842  
...  
ee_007     863  
ee_008     863  
ee_009     863  
ef_000    3486  
eg_000    3485  
Length: 171, dtype: int64
```

```
In [15]: # checking the percentage of missing Values  
df.isnull().mean()*100
```

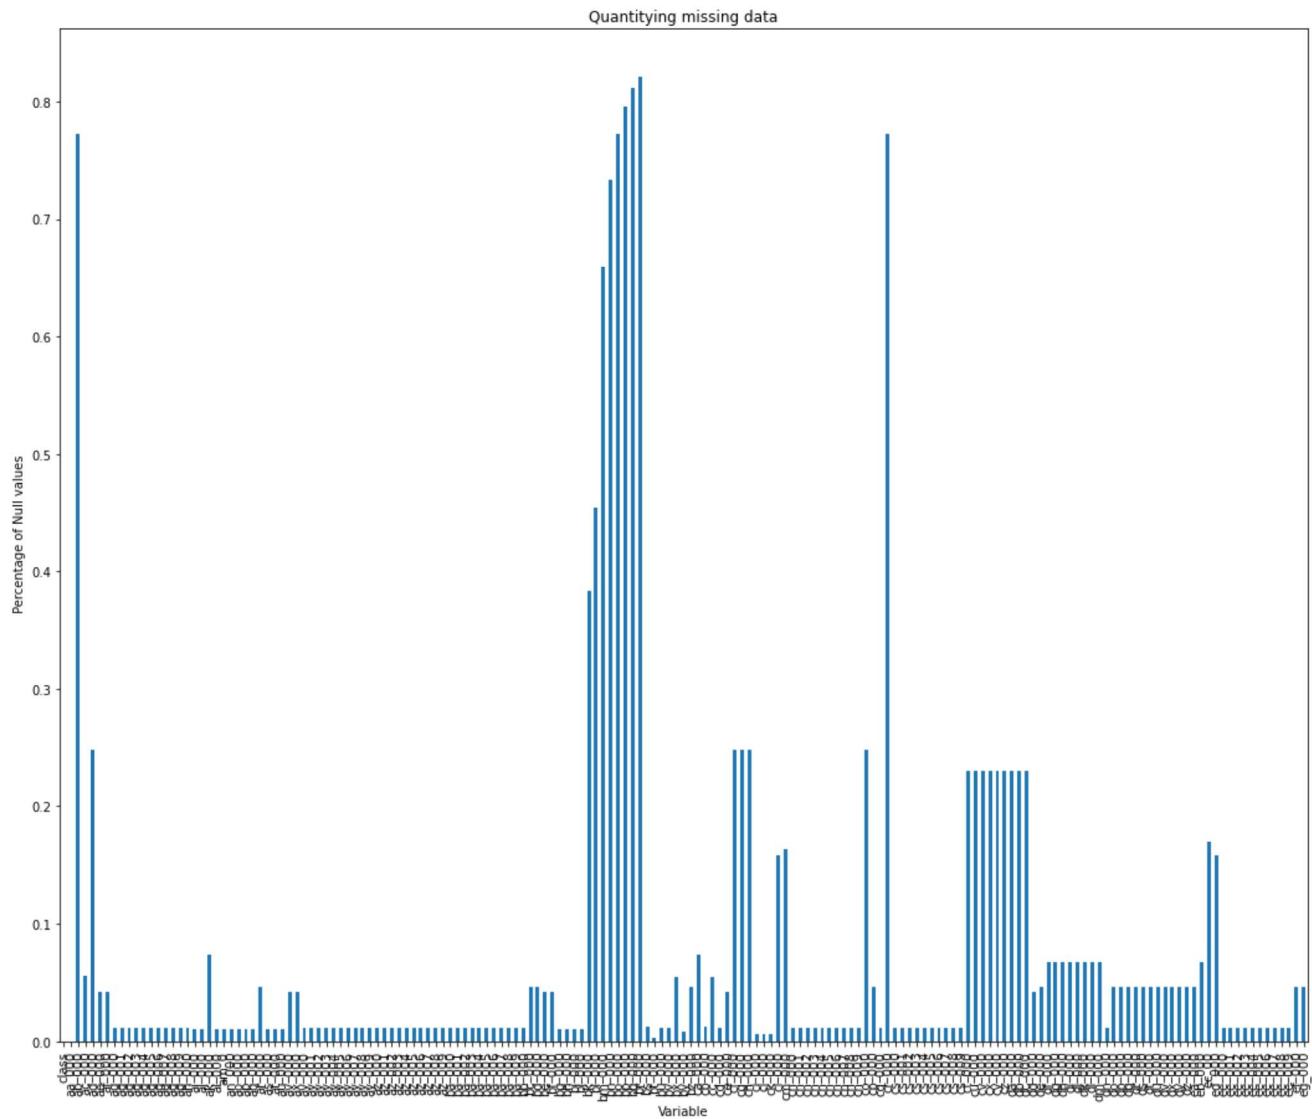
```
Out[15]: class      0.000000  
aa_000      0.000000  
ab_000    77.226316  
ac_000     5.606579  
ad_000    24.792105  
...  
ee_007     1.135526  
ee_008     1.135526  
ee_009     1.135526  
ef_000     4.586842  
eg_000     4.585526  
Length: 171, dtype: float64
```

```
In [16]: df.isnull().sum().sum()
```

```
Out[16]: 1078695
```

```
In [17]: df.isnull().mean().plot.bar(figsize=(18,15))
plt.ylabel(' Percentage of Null values')
plt.xlabel("Variable")
plt.title('Quantitying missing data')
```

```
Out[17]: Text(0.5, 1.0, 'Quantitying missing data')
```



```
In [19]: # first 20 features with highest null values
pd.DataFrame(df.isnull().sum().sort_values(ascending=False)[:10]).rename(columns={0:'Null Value Count'})
```

```
Out[19]:
```

Null Value Count

	Null Value Count
br_000	62393
bq_000	61703
bp_000	60461
bo_000	58709
cr_000	58692
ab_000	58692
bn_000	55722
bm_000	50095
bl_000	34503
bk_000	29128

```
In [21]: # filling null values in feature with median of that feature
for feature in [feature for feature in df.columns if feature not in ['class']]:
    df[feature]=df[feature].fillna(df[feature].median())
```

```
In [23]: ## Checking Null values from top
pd.DataFrame(df.isnull().sum().sort_values(ascending=False)[:5]).rename(columns={0:'Null Value Count'})
```

Out[23]:

Null Value Count	
class	0
cs_003	0
cn_009	0
co_000	0
cp_000	0

In []:

In []:

```
In [24]: # check unique values
df.nunique()
```

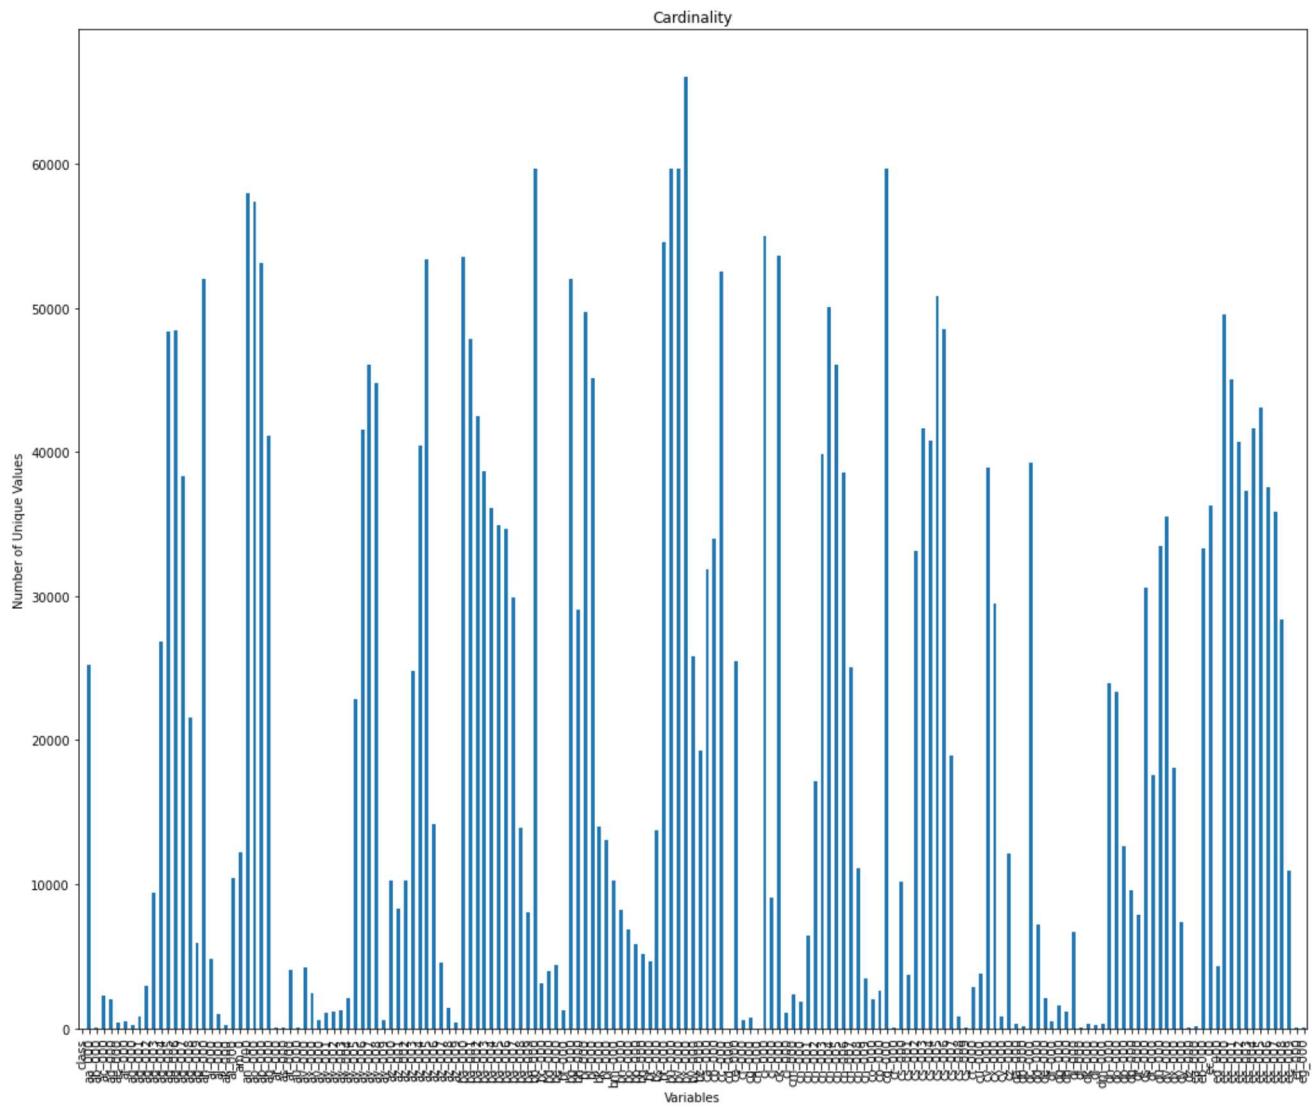
```
Out[24]: class          2
aa_000      25211
ab_000        31
ac_000      2230
ad_000      2028
...
ee_007      35860
ee_008      28413
ee_009      10948
ef_000        31
eg_000        56
Length: 171, dtype: int64
```

```
In [25]: df.nunique().sum()
```

```
Out[25]: 3450493
```

```
In [26]: df.nunique().plot.bar(figsize=(18,15))
plt.ylabel('Number of Unique Values')
plt.xlabel('Variables')
plt.title("Cardinality")
```

```
Out[26]: Text(0.5, 1.0, 'Cardinality')
```



```
In [27]: # check the duplicated values
df.duplicated().sum()
```

```
Out[27]: 0
```

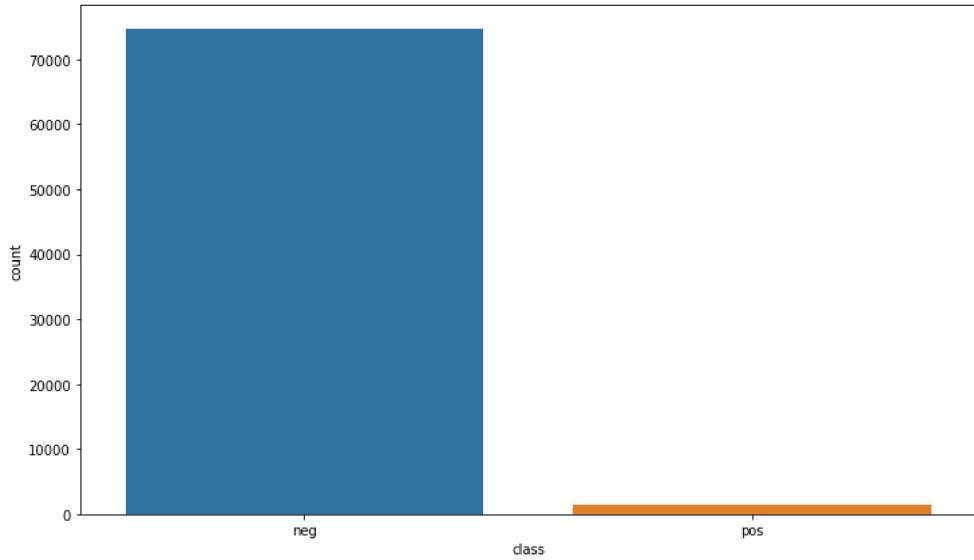
```
In [28]: # check the unique values
df['class'].unique()
```

```
Out[28]: array(['neg', 'pos'], dtype=object)
```

```
In [29]: df['class'].value_counts()
```

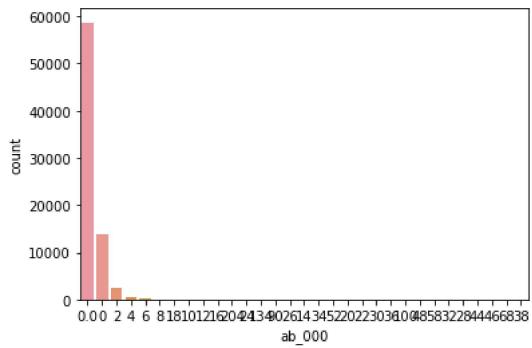
```
Out[29]: neg    74625
          pos    1375
Name: class, dtype: int64
```

```
In [30]: plt.figure(figsize=(12,7))
sns.countplot(df['class'])
plt.show()
```



```
In [31]: sns.countplot(df['ab_000'])
```

```
Out[31]: <AxesSubplot:xlabel='ab_000', ylabel='count'>
```

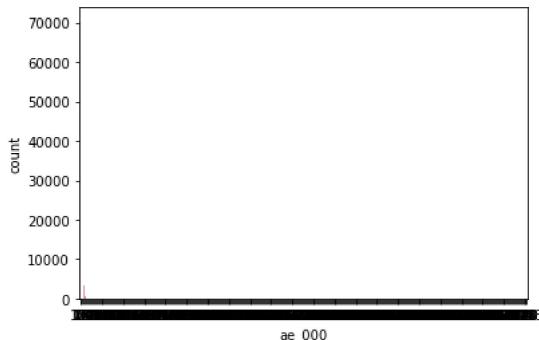


```
In [32]: df['ae_000'].unique()
```

```
Out[32]: array(['0', '16', '104', '0.0', '2', '222', '4', '290', '1286', '170', '6',  
    '806', '1464', '274', '18', '8', '342', '64', '34', '192', '144',  
    '176', '1512', '120', '58', '114', '460', '12', '112', '1314',  
    '20', '550', '36', '1424', '180', '324', '204', '254', '106',  
    '154', '92', '256', '202', '352', '26', '124', '2240', '466',  
    '778', '28', '100', '14', '276', '42', '82', '22', '384', '282',  
    '386', '98', '164', '108', '86', '38', '266', '1056', '134', '10',  
    '200', '150', '326', '350', '398', '424', '32', '128', '752',  
    '314', '732', '2312', '400', '158', '668', '30', '212', '224',  
    '950', '4726', '60', '226', '388', '640', '258', '62', '206', '44',  
    '414', '458', '78', '492', '308', '52', '50', '162', '24', '662',  
    '536', '456', '130', '140', '184', '368', '146', '118', '138',  
    '872', '530', '1118', '356', '218', '126', '102', '1130', '12048',  
    '1508', '396', '374', '658', '816', '556', '216', '66', '760',  
    '136', '220', '280', '46', '462', '788', '368', '416', '348', '74',  
    '1124', '88', '242', '580', '84', '378', '94', '406', '442',  
    '2750', '4698', '40', '748', '292', '952', '284', '270', '1346',  
    '506', '422', '322', '642', '48', '714', '318', '54', '68', '156',  
    '70', '3702', '1654', '454', '1452', '76', '380', '72', '294',  
    '382', '214', '328', '2276', '248', '680', '182', '160', '938',  
    '390', '56', '814', '624', '1200', '1466', '148', '236', '744',  
    '172', '1274', '836', '1582', '804', '2672', '470', '654', '794',  
    '1372', '664', '484', '1468', '168', '110', '1912', '116', '524',  
    '234', '1270', '310', '2468', '612', '228', '486', '246', '2600',  
    '142', '1710', '478', '570', '96', '746', '90', '480', '3478',  
    '538', '1292', '122', '344', '2572', '600', '468', '962', '540',  
    '436', '826', '920', '302', '618', '614', '666', '2322', '2662',  
    '152', '1420', '296', '366', '188', '488', '894', '1714', '572',  
    '522', '702', '320', '208', '238', '500', '762', '542', '336',  
    '446', '21050', '404', '698', '2702', '2306', '548', '1614', '864',  
    '300', '438', '490', '482', '1520', '186', '304', '940', '132',  
    '878', '1012', '376', '2218', '80', '8332', '260', '194', '566',  
    '1852', '4520', '1026', '716', '340', '1860', '166', '1176', '784',  
    '1116', '504', '174', '264', '430', '604', '1398', '2110', '288',  
    '764', '412', '286', '602', '5386', '1412', '2048', '936', '730',  
    '2370', '832', '250', '272', '464', '946', '626', '606', '362',  
    '700', '1700', '402', '782', '526', '516', '1172', '734', '630',  
    '11044', '210', '692', '1502', '262', '244', '972', '230', '2120',  
    '4716', '514', '428', '178', '2088'], dtype=object)
```

```
In [33]: sns.countplot(df['ae_000'])
```

```
Out[33]: <AxesSubplot:xlabel='ae_000', ylabel='count'>
```



```
In [34]: for feature in df.columns:  
    print(df[feature].value_counts())
```

```
neg    74625  
pos    1375  
Name: class, dtype: int64  
8      1293  
10     1201  
12     1158  
14     1135  
6      745  
...  
60302      1  
119474      1  
3088       1  
75834       1  
81852       1  
Name: aa_000, Length: 25211, dtype: int64  
0.0      58692  
0        13842  
2        2484  
4        537  
..     ...
```

Converting Datatype

```
In [37]: # Converting datatypes of numerical feature to int or float from object  
for feature in [feature for feature in df.columns if feature not in ['class', 'aa_000']]:  
    try:  
        df[feature]=df[feature].astype('int64')  
    except:  
        df[feature]=df[feature].astype('float64')
```

```
In [40]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 76000 entries, 0 to 75999  
Columns: 171 entries, class to eg_000  
dtypes: float64(6), int64(164), object(1)  
memory usage: 99.2+ MB
```

```
In [41]: df.dtypes
```

```
Out[41]: class    object  
aa_000    int64  
ab_000    int64  
ac_000    int64  
ad_000    int64  
...  
ee_007    int64  
ee_008    int64  
ee_009    int64  
ef_000    int64  
eg_000    int64  
Length: 171, dtype: object
```

Segregate Numerical and Categorical Features

```
In [42]: # Numerical columns
num_columns=[feature for feature in df.columns if df[feature].dtypes!='object']
num_columns
```

```
Out[42]: ['aa_000',
 'ab_000',
 'ac_000',
 'ad_000',
 'ae_000',
 'af_000',
 'ag_000',
 'ag_001',
 'ag_002',
 'ag_003',
 'ag_004',
 'ag_005',
 'ag_006',
 'ag_007',
 'ag_008',
 'ag_009',
 'ah_000',
 'ai_000',
 'aj_000',
 ...]
```

```
In [43]: len(num_columns)
```

```
Out[43]: 170
```

```
In [44]: # check the Categorical columns
cat_columns=[feature for feature in df.columns if df[feature].dtypes=='object']
cat_columns
```

```
Out[44]: ['class']
```

```
In [45]: len(cat_columns)
```

```
Out[45]: 1
```

```
In [46]: # Getting count of each Numerical values from dataframe
for feature in num_columns:
    print(df[feature].value_counts())
    ...
```

```
284      1
22990     1
9324      1
8670      1
13320     1
Name: ag_002, Length: 2952, dtype: int64
0       60196
30      28
22      28
2       27
16      26
...
1326224    1
9369826    1
5399416    1
432658     1
2388752    1
Name: ag_003, Length: 9383, dtype: int64
...
```

```
In [47]: # check the numerical columns unique values
for feature in num_columns:
    print(df[feature].unique())
    ...
[ 2801180  3477820  1040120 ... 39261520  298950 10074490]
[2445.8  2211.76 1018.64 ... 1714.88 4206.98 1546.74]
[ 2712  2334  1020 ... 20496 5066 5120]
[ 965866  664504  262032 ... 1093142 2830166 883900]
[1706908  824154  453378 ... 12290 1370722 983308]
[1240520  421400  277378 ... 484302 3716 632658]
[493384 178864 159812 ... 129744 51562 201836]
[721044 293306 423992 ... 305924 11802 510354]
[469792 245416 409564 ... 282268 38696 373918]
[339156 133654 320746 ... 345866 111826 349840]
[157956 81140 158022 ... 227924 92102 317840]
[ 73224 97576 95128 ... 177996 298810 960024]
[ 0 1500 514 ... 49346 24458 25566]
[ 0 4 2 8 6 84 22 134 26 10 350 320 40 30 20 482 166 12
276 252 18 82 14 236 144 86 362 74 340 28]
[ 0 32 164 54 4 2 10 6 14 12 8 144 68 86
26 20 28 16 34 94 222 36 892 48 72 152 78 910
70 172 46 80 182 24 56 102 430 42 18 162 40 200
1146 44 108 92 22 606 416 88 124 1720 62 66 168]
```

Statistical Based Analysis

```
In [48]: df.describe()
```

Out[48]:

	aa_000	ab_000	ac_000	ad_000	ae_000	af_000	ag_000	ag_001	ag_002	ag_003	...
count	7.600000e+04	76000.000000	7.600000e+04	7.600000e+04	76000.000000	76000.000000	7.600000e+04	7.600000e+04	7.600000e+04	7.600000e+04	...
mean	6.115976e+04	0.165237	3.364557e+08	1.133174e+05	6.454184	10.381474	1.981626e+02	1.191128e+03	9.587595e+03	9.258930e+04	...
std	2.647366e+05	1.609493	7.769769e+08	3.113852e+07	150.197624	197.330569	1.833153e+04	5.070174e+04	1.709273e+05	8.197977e+05	...
min	0.000000e+00	0.000000	0.000000e+00	0.000000e+00	0.000000	0.000000	0.000000e+00	0.000000e+00	0.000000e+00	0.000000e+00	...
25%	8.600000e+02	0.000000	2.000000e+01	4.200000e+01	0.000000	0.000000	0.000000e+00	0.000000e+00	0.000000e+00	0.000000e+00	...
50%	3.081300e+04	0.000000	1.540000e+02	1.280000e+02	0.000000	0.000000	0.000000e+00	0.000000e+00	0.000000e+00	0.000000e+00	...
75%	4.884000e+04	0.000000	8.480000e+02	2.920000e+02	0.000000	0.000000	0.000000e+00	0.000000e+00	0.000000e+00	0.000000e+00	...
max	4.294967e+07	204.000000	2.130707e+09	8.584298e+09	21050.000000	20070.000000	3.376892e+06	1.047252e+07	1.914916e+07	7.305747e+07	...

8 rows × 170 columns

```
In [49]: df.describe().T
```

Out[49]:

	count	mean	std	min	25%	50%	75%	max
aa_000	76000.0	6.115976e+04	2.647366e+05	0.0	860.0	30813.0	48840.0	4.294967e+07
ab_000	76000.0	1.652368e-01	1.609493e+00	0.0	0.0	0.0	0.0	2.040000e+02
ac_000	76000.0	3.364557e+08	7.769769e+08	0.0	20.0	154.0	848.0	2.130707e+09
ad_000	76000.0	1.133174e+05	3.113852e+07	0.0	42.0	128.0	292.0	8.584298e+09
ae_000	76000.0	6.454184e+00	1.501976e+02	0.0	0.0	0.0	0.0	2.105000e+04
...
ee_007	76000.0	3.440830e+05	1.698180e+06	0.0	122.0	41260.0	166632.0	1.195801e+08
ee_008	76000.0	1.383521e+05	4.667835e+05	0.0	0.0	3862.0	136421.5	1.926740e+07
ee_009	76000.0	8.329316e+03	4.848501e+04	0.0	0.0	0.0	1908.0	4.570398e+06
ef_000	76000.0	8.128947e-02	4.082657e+00	0.0	0.0	0.0	0.0	4.820000e+02
eg_000	76000.0	2.090000e-01	9.999687e+00	0.0	0.0	0.0	0.0	1.720000e+03

170 rows × 8 columns

```
In [50]: # check covariance  
df.cov()
```

Out[50]:

	aa_000	ab_000	ac_000	ad_000	ae_000	af_000	ag_000	ag_001	ag_002	ag_003
aa_000	7.008547e+10	4.240848e+03	-7.785737e+12	-6.601167e+09	9.454919e+05	1.751320e+06	5.671167e+07	1.025314e+09	8.832006e+09	6.672947e+10
ab_000	4.240848e+03	2.590468e+00	-3.112164e+06	-1.866857e+04	8.250324e+00	1.688245e+01	-6.083843e+00	5.582198e+02	6.605238e+03	4.966469e+04
ac_000	-7.785737e+12	-3.112164e+06	6.036931e+17	-3.794526e+13	-4.673741e+08	-1.498406e+09	2.147474e+11	-2.811720e+11	-2.998519e+12	-2.991981e+13
ad_000	-6.601167e+09	-1.866857e+04	-3.794526e+13	9.696074e+14	-7.272759e+05	-1.169290e+06	-2.158851e+07	-1.334675e+08	-1.073488e+09	-1.039052e+10
ae_000	9.454919e+05	8.250324e+00	-4.673741e+08	-7.272759e+05	2.255933e+04	2.460005e+04	-1.278992e+03	-7.260846e+03	2.507139e+04	2.324903e+06
...
ee_007	1.604761e+11	2.937441e+04	-6.450009e+13	-3.819348e+10	8.714962e+06	1.456917e+07	2.382054e+08	5.854768e+09	5.281512e+10	6.621189e+11
ee_008	3.208323e+10	-1.977318e+04	-1.743145e+12	-1.540448e+10	-6.672840e+05	-1.126033e+06	8.762079e+05	1.269533e+08	7.294246e+08	8.006854e+09
ee_009	1.717292e+09	-1.331063e+03	4.527769e+11	-9.391006e+08	-5.338144e+04	-8.615025e+04	-9.614317e+05	-8.678888e+06	-7.512454e+07	-7.158349e+08
ef_000	6.104712e+03	3.399409e-01	-1.966868e+07	-9.147151e+03	3.901344e+01	7.534014e+01	-6.472093e+00	1.986413e+02	3.572187e+03	3.378395e+04
eg_000	2.445737e+04	1.243482e+00	1.264366e+08	-2.354522e+04	4.639863e+01	8.764537e+01	2.648826e+01	1.756286e+03	2.201690e+04	7.879793e+04

170 rows × 170 columns



```
In [51]: # check quantile  
df.quantile()
```

```
Out[51]: aa_000    30813.0  
ab_000      0.0  
ac_000    154.0  
ad_000    128.0  
ae_000      0.0  
...  
ee_007    41260.0  
ee_008    3862.0  
ee_009      0.0  
ef_000      0.0  
eg_000      0.0  
Name: 0.5, Length: 170, dtype: float64
```

```
In [52]: # check count values  
df.count()
```

```
Out[52]: class      76000  
aa_000     76000  
ab_000     76000  
ac_000     76000  
ad_000     76000  
...  
ee_007     76000  
ee_008     76000  
ee_009     76000  
ef_000     76000  
eg_000     76000  
Length: 171, dtype: int64
```

```
In [53]: # check minimum values  
df.min()
```

```
Out[53]: class      neg  
aa_000      0  
ab_000      0  
ac_000      0  
ad_000      0  
...  
ee_007      0  
ee_008      0  
ee_009      0  
ef_000      0  
eg_000      0  
Length: 171, dtype: object
```

```
In [54]: # check maximum values  
df.max()
```

```
Out[54]: class      pos  
aa_000    42949672  
ab_000     204  
ac_000   2130706796  
ad_000   8584297742  
...  
ee_007   119580108  
ee_008   19267396  
ee_009    4570398  
ef_000     482  
eg_000    1720  
Length: 171, dtype: object
```

```
In [55]: # check standard deviation  
df.std()
```

```
Out[55]: aa_000    2.647366e+05  
ab_000    1.609493e+00  
ac_000    7.769769e+08  
ad_000    3.113852e+07  
ae_000    1.501976e+02  
...  
ee_007    1.698180e+06  
ee_008    4.667835e+05  
ee_009    4.848501e+04  
ef_000    4.082657e+00  
eg_000    9.999687e+00  
Length: 170, dtype: float64
```

```
In [56]: # check skewness  
df.skew()
```

```
Out[56]: aa_000    112.926238  
ab_000     53.888614  
ac_000     1.876290  
ad_000    275.680972  
ae_000    91.301964  
...  
ee_007    15.078929  
ee_008    14.354372  
ee_009    33.936322  
ef_000    77.188722  
eg_000    113.494491  
Length: 170, dtype: float64
```

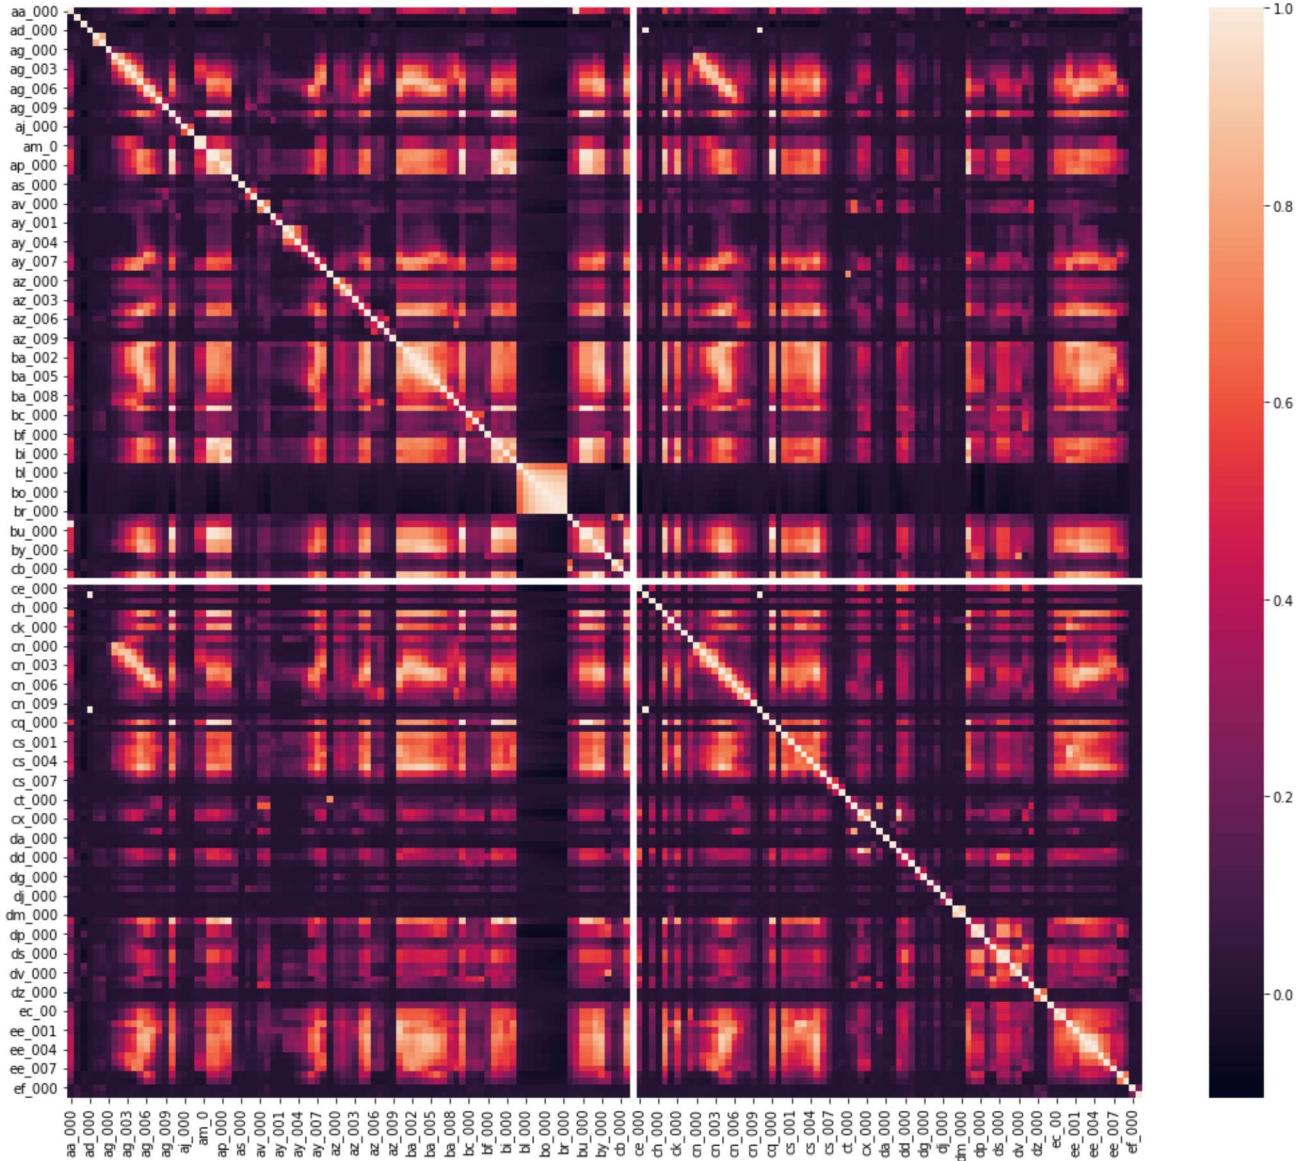
```
In [57]: # check the correlation  
df.corr()
```

```
Out[57]:
```

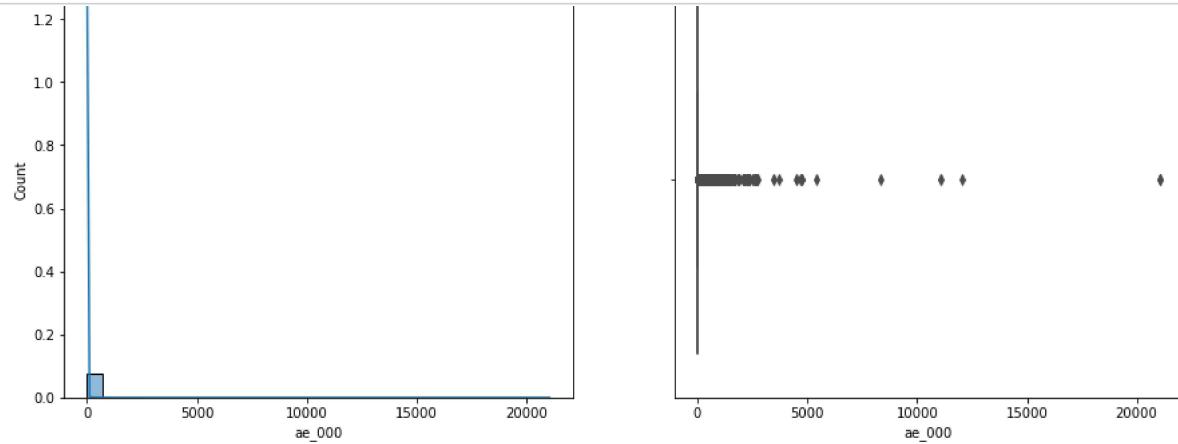
	aa_000	ab_000	ac_000	ad_000	ae_000	af_000	ag_000	ag_001	ag_002	ag_003	...	ee_002	ee_003	ee_004	ee_005
aa_000	1.000000	0.009953	-0.037851	-0.000801	0.023778	0.033524	0.011686	0.076387	0.195179	0.307466	...	0.414098	0.419884	0.409628	0.419756
ab_000	0.009953	1.000000	-0.002489	-0.000372	0.034129	0.053156	-0.000206	0.006841	0.024010	0.037640	...	-0.000332	0.010885	0.001855	0.012075
ac_000	-0.037851	-0.002489	1.000000	-0.001568	-0.004005	-0.009773	0.015077	-0.007137	-0.022578	-0.046973	...	-0.047540	-0.041136	-0.040610	-0.041436
ad_000	-0.000801	-0.000372	-0.001568	1.000000	-0.000156	-0.000190	-0.000038	-0.000085	-0.000202	-0.000407	...	-0.001343	-0.001350	-0.001318	-0.001092
ae_000	0.023778	0.034129	-0.004005	-0.000156	1.000000	0.830001	-0.000465	-0.000953	0.000977	0.018881	...	0.009860	0.016805	0.010357	0.023875
...
ee_007	0.356954	0.010747	-0.048884	-0.000722	0.034168	0.043477	0.007652	0.067999	0.181955	0.475604	...	0.442405	0.433122	0.419859	0.380362
ee_008	0.259626	-0.026319	-0.004806	-0.001060	-0.009518	-0.012225	0.000102	0.005364	0.009142	0.020924	...	0.442837	0.470860	0.448713	0.394357
ee_009	0.133790	-0.017057	0.012019	-0.000622	-0.007330	-0.009004	-0.001082	-0.003530	-0.009065	-0.018009	...	0.224864	0.236117	0.221077	0.184136
ef_000	0.005648	0.051733	-0.006200	-0.000072	0.063622	0.093517	-0.000086	0.000960	0.005119	0.010094	...	0.002271	0.008862	0.001752	0.025645
eg_000	0.009239	0.077262	0.016273	-0.000076	0.030893	0.044417	0.000145	0.003464	0.012881	0.009612	...	0.002475	0.005299	0.005932	0.007154

170 rows × 170 columns

```
In [58]: plt.figure(figsize=(18,15))
sns.heatmap(df.corr())
plt.show()
```

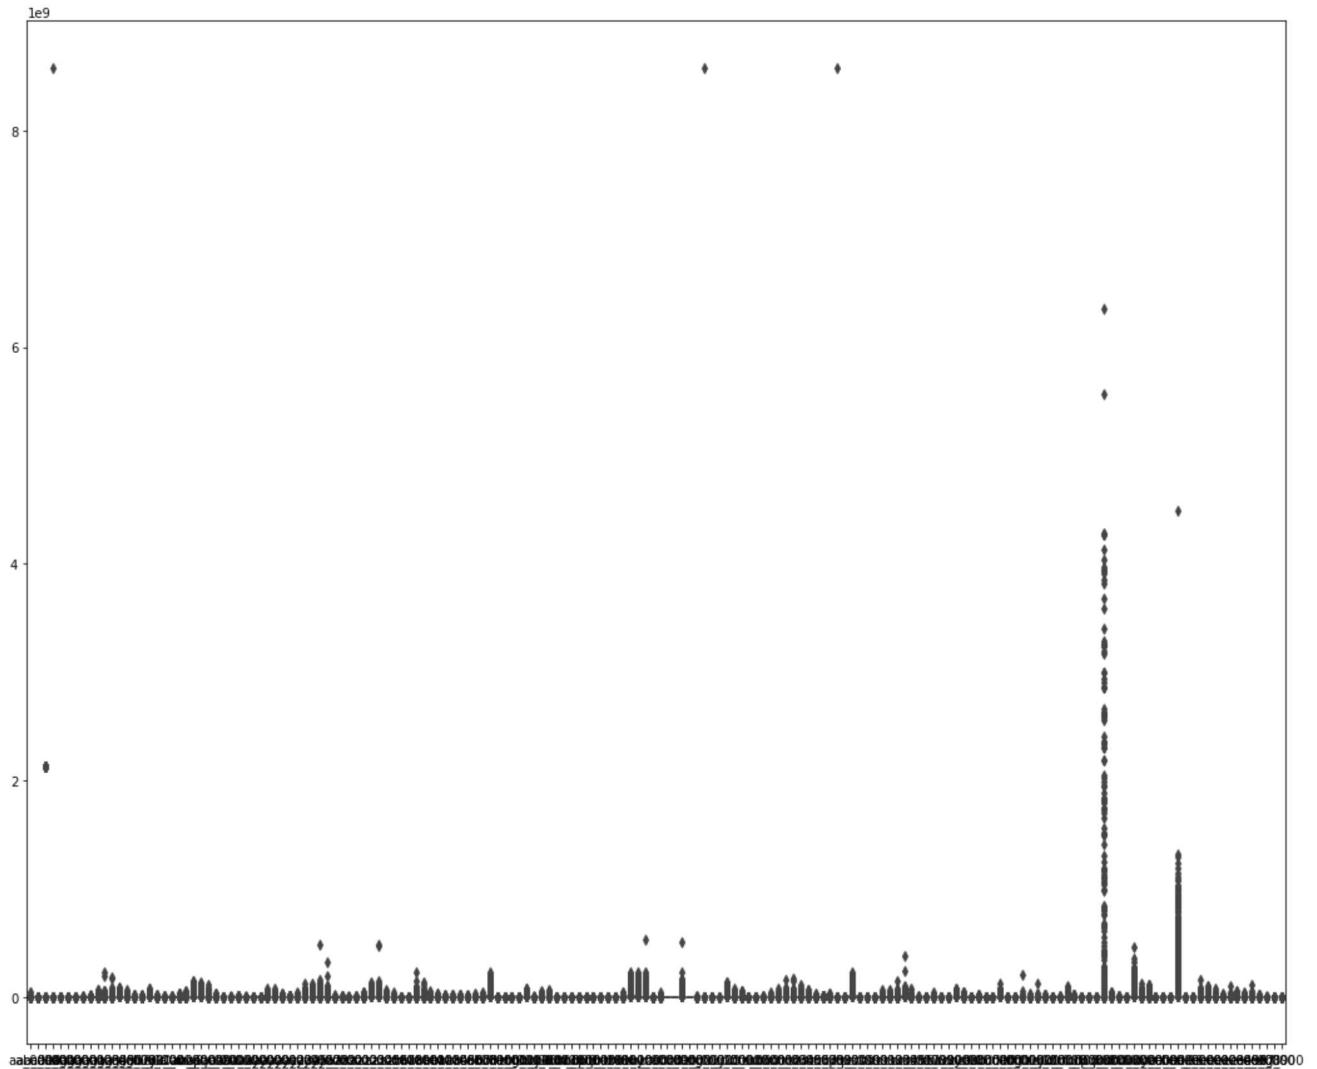


```
In [59]: for feature in num_columns:  
    plt.figure(figsize=(15,6))  
    plt.subplot(121)  
    sns.histplot(data=df, x=feature, kde=True, bins=30)  
    plt.title("{}'s distribution".format(feature), fontweight="bold", fontsize=15)  
  
    plt.subplot(122)  
    sns.boxplot(data=df, x=feature, color='blue')  
    plt.title("{}'s BoxPlot".format(feature), fontweight="bold", fontsize=15)  
    plt.show()
```



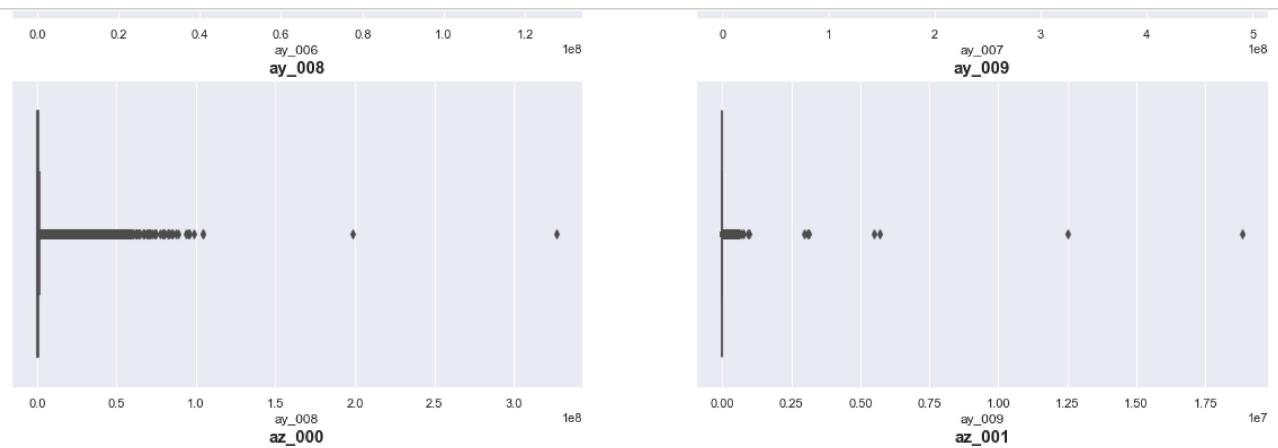
Boxplot

```
In [60]: plt.figure(figsize=(18,15))
sns.boxplot(data=df , orient='v')
plt.show()
```



```
In [61]: ### Checking outliers in numerical features
```

```
plt.figure(figsize=(20,510))
for feature in enumerate(num_columns):
    plt.subplot(85, 2, feature[0]+1)
    sns.set(rc={'figure.figsize':(10,6)})
    sns.boxplot(data=df, x=feature[1], color='red')
    plt.title("{}".format(feature[1]), fontweight="bold", fontsize=15)
```



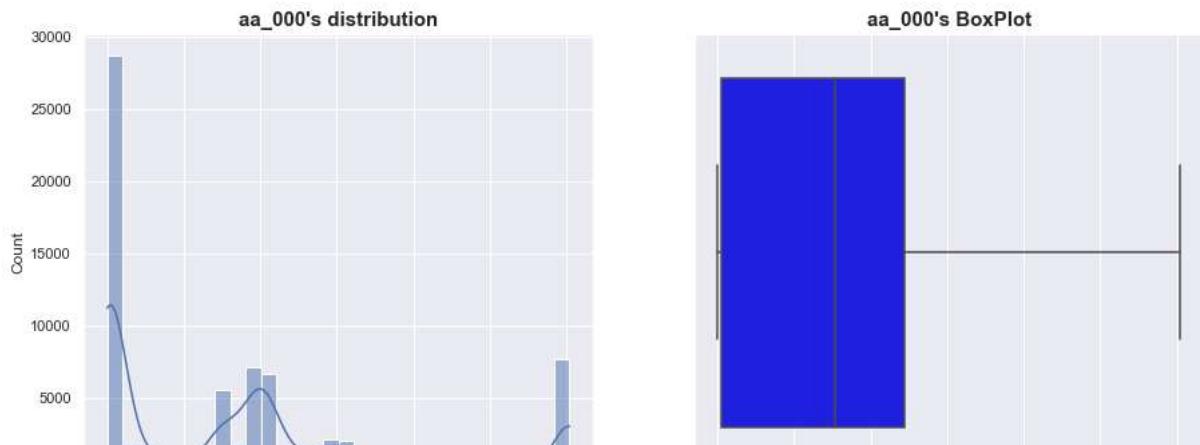
Handle Outliers

```
In [62]: ### Function to Cap Outlier
def remove_outliers(in_data, in_col):
    # Finding the IQR
    first_quartile = in_data[in_col].quantile(0.25)
    third_quartile = in_data[in_col].quantile(0.75)
    iqr = third_quartile - first_quartile
    upper_limit = third_quartile + 1.5 * iqr
    lower_limit = first_quartile - 1.5 * iqr
    in_data.loc[(in_data[in_col]>upper_limit), in_col]= upper_limit
    in_data.loc[(in_data[in_col]<lower_limit), in_col]= lower_limit
    return in_data
```

```
In [63]: ### Capping outliers from continuous features
for feature in num_columns:
    dataset=remove_outliers(df, feature)
```

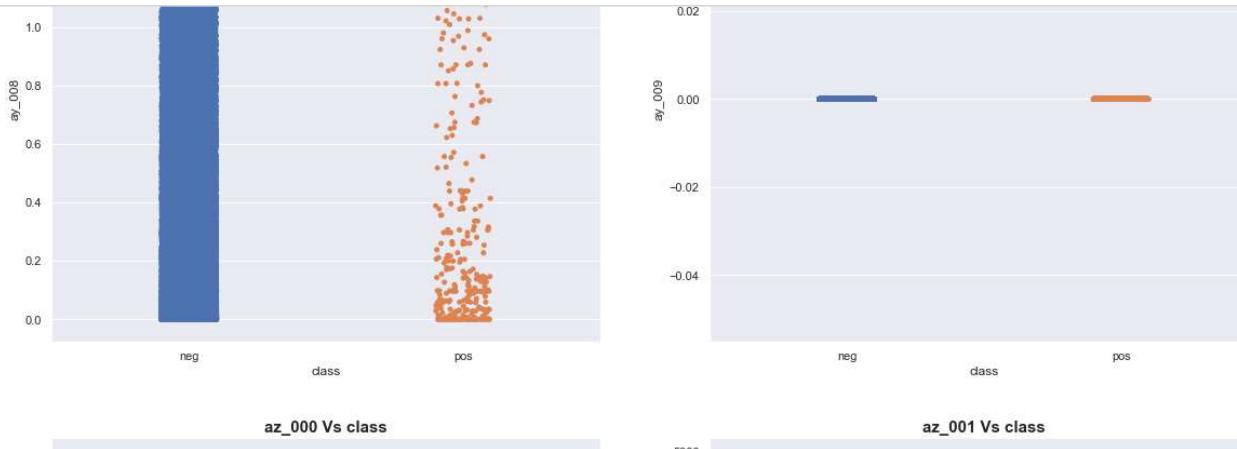
```
In [64]: for feature in num_columns:
    plt.figure(figsize=(15,6))
    plt.subplot(121)
    sns.histplot(data=df, x=feature, kde=True, bins=30)
    plt.title("{}'s distribution".format(feature),fontweight="bold", fontsize=15)

    plt.subplot(122)
    sns.boxplot(data=df, x=feature, color='blue')
    plt.title("{}'s BoxPlot".format(feature),fontweight="bold", fontsize=15)
    plt.show()
```



Numerical Feature vS Target Feature

```
In [184]: ### Checking distribution of Numerical features with respect to target feature class
plt.figure(figsize=(20,840))
for feature in enumerate(num_columns):
    plt.subplot(85, 2, feature[0]+1)
    sns.set(rc={'figure.figsize':(8,10)})
    sns.stripplot(data=df, y=feature[1], x='class')
    plt.title("{} Vs class".format(feature[1]), fontsize=15, fontweight='bold')
```



Dependent and Independent Feature

```
In [67]: # Dependent Feature
X=df.drop('class',axis=1)
X
```

Out[67]:

	aa_000	ab_000	ac_000	ad_000	ae_000	af_000	ag_000	ag_001	ag_002	ag_003	...	ee_002	ee_003	ee_004	ee_005	ee_006	ee_007	ee_008
0	76698	0	2090	280	0	0	0	0	0	0	...	1083002	493384	721044.0	469792	339156.0	157956	73224.00
1	33058	0	0	128	0	0	0	0	0	0	...	421400	178064	293306.0	245416	133654.0	81140	97576.00
2	41040	0	228	100	0	0	0	0	0	0	...	277378	159812	423992.0	409564	320746.0	158022	95128.00
3	12	0	70	66	0	0	0	0	0	0	...	240	46	58.0	44	10.0	0	0.00
4	60874	0	1368	458	0	0	0	0	0	0	...	622012	229790	405298.0	347188	286954.0	311560	341053.75
...	
75995	81852	0	2090	667	0	0	0	0	0	0	...	632658	273242	510354.0	373918	349840.0	317840	341053.75
75996	18	0	52	46	0	0	0	0	0	0	...	266	44	46.0	14	2.0	0	0.00
75997	79636	0	1670	667	0	0	0	0	0	0	...	806832	449962	778826.0	581558	375498.0	222866	341053.75
75998	110	0	36	32	0	0	0	0	0	0	...	588	210	180.0	544	1004.0	1338	74.00
75999	8	0	6	4	0	0	0	0	0	0	...	46	10	48.0	14	42.0	46	0.00

76000 rows × 170 columns

In []:

```
In [68]: # Independent Feature
y=df['class']
y
```

```
Out[68]: 0      neg
1      neg
2      neg
3      neg
4      neg
...
75995    neg
75996    neg
75997    neg
75998    neg
75999    neg
Name: class, Length: 76000, dtype: object
```

```
In [69]: y.head()  
Out[69]: 0    neg  
1    neg  
2    neg  
3    neg  
4    neg  
Name: class, dtype: object
```

```
In [70]: y.unique()  
Out[70]: array(['neg', 'pos'], dtype=object)
```

```
In [71]: ### encoding target feature  
y=y.replace({'neg':0, 'pos':1})
```

```
In [72]: y.unique()  
Out[72]: array([0, 1], dtype=int64)
```

Split_Train_Test

```
In [76]: from sklearn.model_selection import train_test_split
```

```
In [78]: X_train,X_test,y_train,y_test=train_test_split(X,y,test_size=0.35 , random_state=10)
```

```
In [79]: X_train.shape , y_train.shape  
Out[79]: ((49400, 170), (49400,))
```

```
In [81]: X_test.shape , y_test.shape  
Out[81]: ((26600, 170), (26600,))
```

Model Traning

```
In [91]: from sklearn.preprocessing import StandardScaler , MinMaxScaler, RobustScaler  
from sklearn.linear_model import LogisticRegression  
from sklearn.svm import SVC  
from sklearn.tree import DecisionTreeClassifier  
from sklearn.ensemble import RandomForestClassifier, AdaBoostClassifier , GradientBoostingClassifier  
from xgboost import XGBClassifier  
import xgboost  
  
#from sklearn.metrics import accuracy_score , classification_report , ConfusionMatrix  
from sklearn.metrics import precision_score , recall_score , f1_score , roc_auc_score  
  
from sklearn.model_selection import train_test_split , RepeatedStratifiedKFold  
  
from sklearn.compose import ColumnTransformer
```

```
In [92]: # Logistic Regression  
from sklearn.linear_model import LogisticRegression  
logistic=LogisticRegression()  
logistic.fit(X_train , y_train)  
y_pred=logistic.predict(X_test)  
  
from sklearn.metrics import accuracy_score  
accuracy_score_logistic=accuracy_score(y_test , y_pred)  
print(f"Accuracy Score of Logistic Regression Model ={accuracy_score_logistic*100} %")
```

Accuracy Score of Logistic Regression Model =98.28571428571429 %

```
In [94]: # confusion Matrixx  
from sklearn.metrics import confusion_matrix  
conf_mat_log=confusion_matrix(y_test , y_pred)  
print(conf_mat_log)  
tp,tn,fp,fn=conf_mat_log.ravel()
```

[[26046 64] [392 98]]

```
In [96]: conf_mat_accuracy=(tp+tn)/(tp+fp+fn+tn)

Precision=tp/(tp+fp)
Recall = tp/(tp+fn)
f1_score = 2*(Precision * Recall)/(Precision + Recall)
Error_rate=(fp+fn) / (tp+tn+fp+fn)
print(f"Accuracy score confusion matrixx ={round(conf_mat_accuracy * 100,2)} %")
print(f" f1 Score = {f1_score}")
print(f"Precision ={Precision}")
print(f" Recall ={Recall}")
print(f" Misclassification Rate = {Error_rate}")

Accuracy score confusion matrixx =98.16 %
f1 Score = 0.9906812217108516
Precision =0.9851728572509267
Recall =0.9962515299877601
Misclassification Rate = 0.018421052631578946
```

2. Model Training by using Support Vector Classifier

```
In [99]: from sklearn.svm import SVC
sv_class=SVC()
sv_class.fit(X_train, y_train)
y_pred_sv_class=sv_class.predict(X_test)

from sklearn.metrics import accuracy_score
accuracy_score_svc=accuracy_score(y_test , y_pred_sv_class)
print(f"Accuracy Sore of Support Vector Classification ={accuracy_score_svc}")

Accuracy Sore of Support Vector Classification =0.9816165413533835
```

```
In [101]: # Confusion Matrics
from sklearn.metrics import confusion_matrix
conf_mat_svc=confusion_matrix(y_test,y_pred_sv_class)
print(conf_mat_svc)
tp,tn,fp,fn=conf_mat_svc.ravel()

[[26110      0]
 [ 489      1]]
```

```
In [103]: conf_mat_accuracy=(tp+tn)/(tp+tn+fp+fn)
Precision = tp/(tp+fp)
Recall = tp/(tp+fn)
f1_score=2*(Precision * Recall) / (Precision+Recall)
Error_rate= (fp+fn) / (tp+fp+tn+fn)
print(f"Accuracy using Confusion matrix ={round(conf_mat_accuracy*100,2)} %")
print(f"Recall Score = {f1_score}")
print(f"Precision = {Precision}")
print(f"F1 Score ={f1_score}")
print(f"Error Rate ={Error_rate}")
print(f"Misclassification Rate ={Error_rate}")

Accuracy using Confusion matrix =98.16 %
Recall Score = 0.9907038512616202
Precision = 0.981615850219933
F1 Score =0.9907038512616202
Error Rate =0.018421052631578946
Misclassification Rate =0.018421052631578946
```

3. MOdel Training by using Decision Tree Classifier

```
In [105]: from sklearn.tree import DecisionTreeClassifier

model_tree=DecisionTreeClassifier()
model_tree.fit(X_train,y_train)
y_pred_tree=model_tree.predict(X_test)

accuracy_score_tree=accuracy_score(y_test,y_pred_tree)
print(f" Accuracy Score of Decision Tree Classifier ={accuracy_score_tree}")

Accuracy Score of Decision Tree Classifier =0.9842857142857143
```

```
In [106]: # Confusion Matrics

conf_mat_tree=confusion_matrix(y_test, y_pred_tree)
print(conf_mat_tree)
tp,fp,tn,fn = conf_mat_tree.ravel()

[[25891  219]
 [ 199  291]]
```

```
In [108]: conf_mat_accuracy=(tp+fp)/(tp+fp+tn+fp)

Precision = tp/(tp+fp)
Recall = tp/(tp+fn)
f1_score= 2*(Precision * Recall)/(Precision+ Recall)
Error_rate = (fp+fn)/(tp+fp+tn+fp)
print(f" Accuracy using Confusion Matrix ={ round(conf_mat_accuracy * 100,2)}%")
print(f" Precision ={Precision}")
print(f" Recall ={Recall}")
print(f" F1 Score ={f1_score}")
print(f" Error rate ={ Error_rate}")
print(f" Misclassification Rate ={Error_rate}")

Accuracy using Confusion Matrix =98.42%
Precision =0.9916124090386825
Recall =0.9888854938507371
F1 Score =0.9902470741222367
Error rate =0.01922496984318456
Misclassification Rate =0.01922496984318456
```

4. Model Training by using Random Forest Classifier

```
In [145]: from sklearn.ensemble import RandomForestClassifier
random_class=RandomForestClassifier()

random_class.fit(X_train,y_train)
y_pred_random_class=random_class.predict(X_test)

from sklearn.metrics import accuracy_score
accuracy_score_random_class= accuracy_score(y_test,y_pred_random_class)
print(f" Accuracy score of Randon Forest Classifier ={accuracy_score_random_class}")

Accuracy score of Randon Forest Classifier =0.991015037593985
```

```
In [120]: # Confusion Matrics

conf_mat_random_fore= confusion_matrix(y_test, y_pred_random_class)
print(conf_mat_random_fore)
tp,fp,tn,fn = conf_mat_random_fore.ravel()

[[26084   26]
 [ 216  274]]
```

```
In [150]: conf_mat_accuracy= (tp+tn)/(tp+fp+tn+fn)

Precision=tp/(tp+fp)
Recall=tp/(tp+fn)
f1_score = 2*(Precision*Recall)/(Precision+Recall)
Error_rate= (fp+fn)/(tp+fp+tn+fn)
print(f"Accuracy using confusion matrix = {round(conf_mat_accuracy*100,2)}%")
print(f"Precision ={Precision}")
print(f"Recall = {Recall}")
print(f" F1 Score ={f1_score}")
print(f" Error Rate ={Error_rate}")
print(f" Misclassification Rate={Error_rate}")

Accuracy using confusion matrix = 98.8%
Precision =0.9984680199157411
Recall = 0.9894113628600706
F1 Score = 0.9939190605997065
Error Rate =0.011992481203007518
Misclassification Rate=0.011992481203007518
```

5. Model Training by using Extra tree Classifier

```
In [148]: from sklearn.ensemble import ExtraTreesClassifier
et_model=ExtraTreesClassifier()
et_model.fit(X_train,y_train)
y_pred_et_model=et_model.predict(X_test)

from sklearn.metrics import accuracy_score
accuracy_score_et_model=accuracy_score(y_test,y_pred_et_model)
print(f" Accuracy score of Extra Tree Classiier ={accuracy_score_et_model}")

Accuracy score of Extra Tree Classiier =0.9903383458646616
```

```
In [126]: # Confusion Matrics

conf_mat_et_model = confusion_matrix(y_test,y_pred_et_model)
print(conf_mat_et_model)
tp,fp,tn,fn=conf_mat_et_model.ravel()

[[26070    40]
 [ 211   279]]
```

```
In [149]: conf_mat_accuracy= (tp+tn)/(tp+fp+tn+fn)

Precision=tp/(tp+fp)
Recall=tp/(tp+fn)
f1_score = 2*(Precision*Recall)/(Precision+Recall)
Error_rate= (fp+fn)/(tp+fp+tn+fn)
print(f"Accuracy using Confusion Matrix = {round(conf_mat_accuracy*100,2)}")
print(f"Precision ={Precision}")
print(f"Recall = {Recall}")
print(f"F1 Score = {f1_score}")
print(f"Error Rate ={Error_rate}")
print(f" Misclassification Rate={Error_rate}")

Accuracy using Confusion Matrix = 98.8
Precision =0.9984680199157411
Recall = 0.9894113628600706
F1 Score = 0.9939190605997065
Error Rate =0.011992481203007518
Misclassification Rate=0.011992481203007518
```

6. Model Training by Using Bagging Classifier

```
In [151]: from sklearn.ensemble import BaggingClassifier

bagg_model=BaggingClassifier()
bagg_model.fit(X_train,y_train)
y_pred_bagg_model=bagg_model.predict(X_test)

accuracy_score_bagg_model=accuracy_score(y_test,y_pred_bagg_model)
print(accuracy_score_bagg_model)
print(f"Accuracy Score of Bagging Claasifier ={accuracy_score_bagg_model}")

0.9892481203007519
Accuracy Score of Bagging Claasifier =0.9892481203007519
```

```
In [166]: # Confusion Model

conf_mat_bagg_model = confusion_matrix(y_test,y_pred_bagg_model)
print(conf_mat_bagg_model)
tp,fp,tn,fn=conf_mat_bagg_model.ravel()

[[26054    56]
 [ 230   260]]
```

In []:

```
In [167]: conf_mat_bagg_model= (tp+tn)/(tp+fp+tn+fn)

Precision=tp/(tp+fp)
Recall=tp/(tp+fn)
f1_score = 2*(Precision*Recall)/(Precision+Recall)
Error_rate= (fp+fn)/(tp+fp+tn+fn)
print(f"Accuracy using Confusion Matrix = {round(conf_mat_bagg_model*100)}")
print(f"Precision ={Precision}")
print(f"Recall = {Recall}")
print(f"F1 Score = {f1_score}")
print(f"Error Rate ={Error_rate}")
print(f" Misclassification Rate={Error_rate}")


```

```
Accuracy using Confusion Matrix = 99
Precision =0.9978552278820375
Recall = 0.9901193281143118
F1 Score = 0.9939722264611628
Error Rate =0.0118796992481203
Misclassification Rate=0.0118796992481203
```

7. Model Traning by using AdaBoost Classifier

```
In [158]: from sklearn.ensemble import AdaBoostClassifier
ada_model=AdaBoostClassifier()
ada_model.fit(X_train,y_train)
y_pred_ada=ada_model.predict(X_test)

from sklearn.metrics import accuracy_score
accuracy_score_ada=accuracy_score(y_test,y_pred_ada)
print(f"Accuracy score of Ada BoostClaasifier={accuracy_score_ada}")


```

```
Accuracy score of Ada BoostClaasifier=0.9865413533834586
```

```
In [159]: # confusion matrix

conf_mat_ada=confusion_matrix(y_test,y_pred_ada)
print(conf_mat_ada)
tp,fp,tn,fn=conf_mat_ada.ravel()

[[26007  103]
 [ 255  235]]
```

```
In [162]: conf_mat_accuracy=(tp+tn)/(tp+fp+tn+fn)

Precision = tp/(tp+fp)
Recall=tp/(tp+tn)
f1_score=2*(Precision*Recall)/(Precision+Recall)
Error_rate=(fp+fn)/(tp+fp+tn+fn)
print(f"Accuracy using confusion matrix={round(conf_mat_accuracy*100)}")
print(f"Precision ={Precision}")
print(f"Recall={Recall}")
print(f"F1 Score={f1_score}")
print(f"Error_rate={Error_rate}")


```

```
Accuraccy using confusion matrix=99
Precision =0.9960551512830333
Recall=0.9902901530728809
F1 Score=0.9931642862598336
Error_rate=0.012706766917293232
```

8. Model Trainig using by Gradient Boosting Classifier

```
In [171]: from sklearn.ensemble import GradientBoostingClassifier
boost_model=GradientBoostingClassifier()
boost_model.fit(X_train,y_train)
y_pred_boost_model=boost_model.predict(X_test)

accuracy_score_boost_model=accuracy_score(y_test,y_pred_boost_model)
print(f"Accuracy score  of GradientBoostClassifier={GradientBoostingClassifier}")


```

```
Accuracy score  of GradientBoostClassifier=<class 'sklearn.ensemble._gb.GradientBoostingClassifier'>
```

```
In [182]: # confusion matrix

conf_mat_boost=confusion_matrix(y_test,y_pred_boost_model)
print(conf_mat_boost)
tp,fp,tn,fn=conf_mat_boost.ravel()

[[26047    63]
 [ 228   262]]
```

```
In [183]: conf_mat_accuracy=(tp+fp)/(tp+fp+tn+fn)

Precision=(tp)/(tp+fp)
Recall=(tp)/(tp+fn)
f1_score=2*(Precision*Recall)/(Precision+Recall)
Error_rate=(fp+fn)/(tp+fp+tn+fn)
print("Accuracy Score of confusion matrix={round(conf_mat_accuracy*100)}%")
print(f"Precision={Precision}")
print(f"Recall={Recall}")
print(f"F1 Score={f1_score}")
print(f"Error Rate ={Error_rate}")
print(f"MisClassification Rtae ={Error_rate}")

Accuracy Score of confusion matrix=98%
Precision=0.9975871313672923
Recall=0.9900414306891178
F1 Score=0.9937999580304852
Error Rate =0.012218045112781954
MisClassification Rtae =0.012218045112781954
```

9.Extreme GradientBoostingClassifier Model ---XGB

```
In [176]: import xgboost
from xgboost import XGBClassifier

model_xgb=xgboost.XGBClassifier()
model_xgb.fit(X_train,y_train)
y_pred_xgb=model_xgb.predict(X_test)

accuracy_score_xgb=accuracy_score(y_test,y_pred_xgb)
print(f" Accuracy score of XGB model ={accuracy_score_xgb}")

Accuracy score of XGB model =0.9921428571428571
```

```
In [180]: # confusion matrix

conf_mat_xgb=confusion_matrix(y_test,y_pred_xgb)
print(conf_mat_xgb)
tp,fp,tn,fn=conf_mat_xgb.ravel()

[[26062    48]
 [ 161   329]]
```

```
In [181]: conf_mat_accuracy=(tp+fp)/(tp+fp+tn+fn)

Precision=(tp)/(tp+fp)
Recall=(tp)/(tp+fn)
f1_score=2*(Precision*Recall)/(Precision+Recall)
Error_rate=(fp+fn)/(tp+fp+tn+fn)
print("Accuracy Score of confusion matrix={round(conf_mat_accuracy*100)}%")
print(f"Precision={Precision}")
print(f"Recall={Recall}")
print(f"F1 Score={f1_score}")
print(f"Error Rate ={Error_rate}")
print(f"MisClassification Rtae ={Error_rate}")

Accuracy Score of confusion matrix=98%
Precision=0.9981616238988893
Recall=0.9875336288886363
F1 Score=0.9928191843964876
Error Rate =0.014172932330827068
MisClassification Rtae =0.014172932330827068
```

Final Result of All Classifier Model

```
In [190]: print(f"1. Accuracy Score of Logistic Regression Model ={accuracy_score_logistic*100} %")
print(f"2. Accuracy Score of Support Vector Classification ={accuracy_score_svc*100} %")
print(f"3. Accuracy Score of Decision Tree Classifier ={accuracy_score_tree*100} %")
print(f"4. Accuracy score of Random Forest Classifier ={accuracy_score_random_class*100} %")
print(f"5. Accuracy score of Extra Tree Classifier ={accuracy_score_et_model*100} %")
print(f"6. Accuracy Score of Bagging Classifier ={accuracy_score_bagg_model*100} %")
print(f"7. Accuracy score of Ada BoostClassifier={accuracy_score_ada*100} %")
print(f"8. Accuracy score of GradientBoostClassifier={GradientBoostingClassifier} %")
print(f"9. Accuracy score of XGB model ={accuracy_score_xgb*100} %")
```

1. Accuracy Score of Logistic Regression Model =98.28571428571429 %
2. Accuracy Score of Support Vector Classification =98.16165413533835%
3. Accuracy Score of Decision Tree Classifier =98.42857142857143%
4. Accuracy score of Random Forest Classifier =99.1015037593985%
5. Accuracy score of Extra Tree Classifier =99.03383458646616%
6. Accuracy Score of Bagging Classifier =98.92481203007519 %
7. Accuracy score of Ada BoostClassifier=98.65413533834587 %
8. Accuracy score of GradientBoostClassifier=<class 'sklearn.ensemble._gb.GradientBoostingClassifier'> %
9. Accuracy score of XGB model =99.21428571428571 %

Final Report

Highest Accuracy Score is XGB Classifier 99.2142%

Thank you.....

```
In [ ]:
```